

Determining Language Grounding Failures in Vision-Language-Action Models

Anonymous ACL submission

Abstract

Vision-Language-Action (VLA) models report high aggregate success rates on manipulation benchmarks, but these single-number metrics hide whether models genuinely follow instructions or exploit visual shortcuts and memorized scene layouts. In this paper, we isolate the language channel using a scene-fixed causal evaluation protocol on `LIBERO-Goal`, holding the initial simulator state strictly constant while perturbing instructions across three architectures spanning a $16\times$ parameter range (`SmolVLA-450M`, `OpenVLA-7B`, and `OpenVLA-OFT-7.5B`).

We find that language is causally necessary and scale-invariant: destroying the instruction by blanking or scrambling it collapses task success across all scales. However, we show that widely reported word-class asymmetries—where models appear to bind object nouns strictly while treating action verbs as near-optional—are confounded by benchmark information structure. In `LIBERO-Goal`, the noun set uniquely identifies all ten tasks ($H(\text{task} \mid \text{nouns}) = 0$ bits), rendering the verb redundant. When we test genuine verb antonyms, `OpenVLA-7B` drops 31.7 points ($p < 10^{-5}$), proving that it reads the verb when the benchmark requires it. We present an episode-level goal predicate taxonomy to distinguish instruction-following from scene-memorized defaults, and argue that benchmark information content must be reported before declaring models instruction-blind.

1 Introduction

Vision-Language-Action (VLA) models combine vision-language backbones with robotic control heads, promising policies that execute physical manipulations driven by natural-language instructions (Brohan et al., 2022, 2023; Driess et al., 2023; Kim et al., 2025b; Shukor et al., 2025). Standard evaluation benchmarks report high aggregate

success rates, and the community frequently interprets these figures as evidence that policies genuinely understand their instructions. However, a model can score well on a benchmark without parsing language at all by exploiting visual shortcuts, memorized task layouts, or superficial keyword matching. This gap between aggregate success and genuine language grounding is the central problem we address in this paper.

We approach this problem through two core research questions:

- **RQ1:** Does language causally drive a VLA model’s physical behavior when the initial visual scene is held strictly constant?
- **RQ2:** Which linguistic operations break grounding, and do observed failures stem from object reference (nouns) or action reference (verbs)?

To answer these questions, we hold the physical simulator scene fixed and treat only the instruction text as the experimental variable. This design choice separates our work from existing VLA robustness benchmarks, which vary multiple factors simultaneously—such as object positions, camera angles, lighting, sensor noise, and language—or compound several perturbations at once. By isolating language as the sole manipulated variable, we determine the instruction’s exact role in driving policy behavior rather than making an aggregate robustness claim.

Evaluating three model architectures spanning a $16\times$ parameter range on `LIBERO-Goal`, we find that language is causally load-bearing and scale-invariant: destroying the instruction via blanking or length-matched nonsense collapses task success across all model scales. However, when we investigate word-class sensitivity, we discover that the widely reported asymmetry—where models seem sensitive to object nouns but insensitive to

action verbs—is an artifact of benchmark information structure. In LIBERO-Goal, the noun set carries all 3.32 bits of task identity ($H(\text{task} \mid \text{nouns}) = 0$ bits), rendering the verb fully redundant. When we isolate the three tasks with true verb antonyms, OpenVLA-7B loses 31.7 points ($p < 10^{-5}$), demonstrating that it reads the verb when the benchmark requires it. We propose an episode-level goal-predicate taxonomy and metric suite to separate genuine instruction compliance from scene-memorized defaults.

2 Preliminaries

2.1 Vision-Language-Action Models

The VLA paradigm originated with RT-1 (Brohan et al., 2022), which framed robotic control as a sequence-modeling problem over tokenized actions. RT-2 (Brohan et al., 2023) showed that web-scale vision-language pre-training transfers into control, yielding emergent semantic understanding, while PaLM-E (Driess et al., 2023) demonstrated that a unified multimodal model benefits from joint training across language, vision, and control. Open X-Embodiment (Collaboration et al., 2023) scaled training data across more than a million real trajectories. These foundational works establish an expectation: if a policy inherits a language model’s pre-trained semantics, it should condition its physical actions on the instruction. This expectation is precisely what our study puts to the test.

2.2 Evaluated Models

We evaluate three open-source checkpoints that instantiate the primary architectural choices in current VLA research:

1. **SmolVLA-450M** (Shukor et al., 2025): A compact policy combining a SmolVLM2 backbone with a *flow-matching* action expert, serving as our small-scale baseline.
2. **OpenVLA-7B** (Kim et al., 2025b): The canonical 7B VLA policy built on Prismatic VLM and Llama-2, emitting *discretized* action tokens from a single static camera.
3. **OpenVLA-OFT-7.5B** (Kim et al., 2025a): The Optimized Fine-Tuning recipe utilizing continuous action chunking, an L1-regression loss, an additional wrist camera stream, and 8-dimensional proprioception.

(Note that OFT here refers to *Optimized* Fine-Tuning, not orthogonal fine-tuning; the checkpoint is evaluated without FiLM.)

These three models span a $16\times$ parameter range and three distinct action parameterizations (flow-matching, discretized tokens, and continuous L1 regression), allowing us to test whether language grounding behavior generalizes across architectures.

2.3 Benchmarks

We conduct all evaluations within the LIBERO benchmarking framework (Liu et al., 2023). LIBERO contains four distinct task suites: Spatial, Object, Goal, and Long. We focus exclusively on LIBERO-Goal because its ten tasks share an identical initial scene layout and object inventory. In LIBERO-Goal, the initial visual state S_0 does not disambiguate the goal, making the instruction the sole distinguishing input.

3 Related Work

VLA Robustness and Benchmark Shortcuts.

Recent studies highlight severe vulnerabilities in VLA policies. Pang et al. (2025) evaluate the positional robustness of OpenVLA on LIBERO-Object, finding that success drops sharply outside the specific object positions seen during training, indicating rote memorization. LIBERO-Plus (Fei et al., 2025) benchmarks over 10,000 tasks across seven variation axes, reporting that language perturbations cause minor performance drops while visual perturbations disrupt policies severely. LIBERO-Para (Kim et al., 2026) evaluates paraphrase robustness across object and action references and introduces the PRIDE metric, which we adopt to measure paraphrase sensitivity.

The Instruction-Blindness Consensus.

A growing literature argues that VLAs largely ignore language. LIBERO-PRO (Zhou et al., 2026) reports that outputs remain unchanged under corrupted instructions, attributing success to memorized action sequences. LangGap (Hou and Zhao, 2026) demonstrates that high-scoring models fail when forced to use language. Zhang and Li (2026) term this failure mode “instruction blindness”, while Zhan and Wang (2026) report a “modality collapse” toward visual affordances. Mechanistically, Grant et al. (2026) show that

injecting baseline activations into null-prompt episodes recovers baseline behavior across six models.

Linguistic Shortcuts in NLP. Reliance on superficial dataset shortcuts is a well-documented challenge in computational linguistics. In Natural Language Inference (NLI), models routinely exploit lexical overlap and syntactic heuristics (Gururangan et al., 2018; McCoy et al., 2019). In Visual Question Answering (VQA), models frequently exploit language priors to answer without inspecting image details (Goyal et al., 2017). Our work connects this NLP dataset-artifact literature to VLA evaluation, demonstrating that apparent "verb blindness" in VLAs is an artifact of information redundancy in benchmark instruction design.

4 Language Grounding Failures

4.1 Defining Language Grounding Failure

We define a *language grounding failure* as an episode in which the policy’s physical behavior is not causally driven by the linguistic content of the instruction. We distinguish this from an *execution failure*, in which the policy correctly infers the intended goal but fails to physically execute the trajectory.

Scalar success rates conflate these two independent properties: whether the policy understood the instruction, and whether the policy possessed the motor skill to complete it. A model can fail a task despite perfect language grounding (producing a correct approach trajectory that slips at the grasp phase). Conversely, a model can succeed at a task despite zero language grounding if its behavior is driven by memorized scene layouts or visual priors.

We resolve this ambiguity by holding the initial simulator state S_0 strictly constant while perturbing the instruction I . We then evaluate the final simulator state against all goal predicates achievable in that scene, rather than relying solely on the single predicate associated with the original task.

4.2 The Perturbation Taxonomy

Let I_{orig} denote the canonical instruction for a task on fixed scene S_0 . We construct a family of perturbed instructions I_{cond} grouped by the shortcut hypothesis each condition falsifies:

- **Destructive Conditions:** Remove instruction content entirely while leaving scene S_0

untouched. `blank` supplies an empty string; `nonsense` supplies length-matched out-of-domain tokens. If the policy relies on language, success must collapse. If the policy succeeds, its behavior is driven by scene memorization.

- **Misleading Conditions:** Substitute a valid, executable instruction describing a different goal in S_0 . `wrong_object` replaces the target object noun; `wrong_action` replaces the action verb; `wrong_task` replaces the full instruction with another task’s instruction. These conditions test whether a policy follows the new instruction (compliance) or defaults to the scene’s original goal (memorization).

- **Control & Paraphrase Conditions:** `repeated` concatenates I_{orig} with itself to test sensitivity to length. Paraphrase conditions (`para_object`, `para_action`, `para_compositional`) reword the instruction along specific linguistic axes while preserving meaning, testing robustness to surface form.

Table 1 summarizes the perturbation taxonomy.

4.3 Ground-Truth Evaluation via Goal Predicates

Standard simulators evaluate rollouts using a single boolean flag corresponding to the original task predicate g_{k^*} . Under perturbed instructions, this creates a mismatch: a policy that correctly follows a misleading instruction (e.g., executing `wrong_task`) is mis-scored as a failure by the native simulator check.

We exploit a structural property of LIBERO-Goal: all ten tasks are defined over the identical initial scene S_0 . We evaluate every task’s goal predicate g_k against the final state S_T of every rollout, constructing the *achieved task set*:

$$\mathcal{A}(S_T) = \{k : g_k(S_T) = \text{true}\} \quad (1)$$

where $k \in \{1, \dots, 10\}$. $\mathcal{A}(S_T)$ reveals whether the policy achieved no goal ($\emptyset$), the original goal (k^*), or a newly instructed goal (k_{wrong}).

4.4 Evaluating Perturbed Instructions via Achieved Goal Sets

Using $\mathcal{A}(S_T)$, we establish per-episode classification schemes for destructive and misleading per-

Condition	Construction	Shortcut Hypothesis Ruled Out
original	I_{orig} , unchanged	Baseline reference.
blank	Empty string	Policy succeeds without language (pure scene memorization).
nonsense	Length-matched out-of-domain tokens	Policy relies on keyword presence without structural semantics.
wrong_object	Swapped target object noun	Policy ignores named object and defaults to a fixed target.
wrong_action	Swapped action verb	Policy ignores named action and defaults to a fixed manipulation.
wrong_task	Full instruction of another in-scene task	Policy ignores instruction and executes scene’s default task.
repeated	I_{orig} concatenated with itself	Sensitivity to input length rather than meaning.
para_object, para_action, para_compositional	Reworded instruction preserving original goal	Policy is brittle to surface form despite unchanged semantics.

Table 1: The perturbation taxonomy: instruction construction and the specific shortcut hypothesis each condition rules out.

turbations.

4.4.1 Destructive Conditions

For `blank` and `nonsense`, no valid goal is specified. We cross whether I_{orig} succeeded against whether $\mathcal{A}(S_T) \neq \emptyset$ under the destructive condition (Table 2):

- **Grounded:** I_{orig} succeeded and $\mathcal{A}(S_T) = \emptyset$. Removing language eliminated success.
- **Memorization:** I_{orig} succeeded and $\mathcal{A}(S_T) \neq \emptyset$. The policy completed a task without language.
- **Spurious Success:** I_{orig} failed but $\mathcal{A}(S_T) \neq \emptyset$ under destructive prompt.
- **Uninformative:** I_{orig} failed and $\mathcal{A}(S_T) = \emptyset$. The policy could not complete the task under any prompt.

We define the episode-level Grounding Rate (GR) as:

$$\text{GR}(c) = \frac{|\text{Grounded}|}{|\text{Grounded}| + |\text{Mem}| + |\text{Spurious}|} \quad (2)$$

4.4.2 Misleading Conditions

For `wrong_object`, `wrong_action`, and `wrong_task`, the perturbed prompt specifies goal $k_{\text{wrong}} \neq k^*$. We classify final states into three outcomes (Table 3):

- **Follow:** $k_{\text{wrong}} \in \mathcal{A}(S_T)$. The policy executed the newly instructed task (instruction compliance).

- **Ignore:** $k^* \in \mathcal{A}(S_T)$ and $k_{\text{wrong}} \notin \mathcal{A}(S_T)$. The policy executed the original task despite the altered prompt (scene-prior override).

- **Fail:** Neither goal achieved.

4.5 Metrics

We convert episode classifications into four aggregate metrics:

1. **Causal Sensitivity Score (CSS):** For destructive conditions, $\text{CSS}(c) = \frac{\text{TSR}_0 - \text{TSR}(c)}{\max(\text{TSR}_0, \epsilon)}$. $\text{CSS} \approx 1$ indicates strong causal necessity; $\text{CSS} \approx 0$ indicates language is ignored.
2. **Original-Action Rate (OAR):** For misleading conditions, $\text{OAR}(c) = \frac{|\text{Ignore}|}{|\text{Follow}| + |\text{Ignore}| + |\text{Fail}|}$. High OAR measures scene-memorization overrides.
3. **PRIDE Score:** For paraphrases, we compute task success weighted by linguistic paraphrase distance $\text{PD} = 1 - 0.5(\text{SK} + \text{ST})$, where SK and ST measure keyword and structural similarity (Kim et al., 2026).
4. **Failure Locus:** Failed episodes are categorized into *planning failures* (end-effector never approaches target within distance threshold τ) versus *execution failures* (end-effector reaches target but fails manipulation).

5 Experiments

5.1 Models and Benchmark

We evaluate SmolVLA-450M, OpenVLA-7B, and OpenVLA-OFT-7.5B on the ten tasks of

	No Goal Achieved ($\mathcal{A}(S_T) = \emptyset$)	Some Goal Achieved ($\mathcal{A}(S_T) \neq \emptyset$)
I_{orig} Succeeded	Grounded. Removing language eliminated task success.	Memorization. Policy succeeded without language; driven by scene priors.
I_{orig} Failed	Uninformative. Policy incompetent on this scene; excluded from grounding rate.	Spurious Success. Policy achieved a goal without language where original prompt failed.

Table 2: The destructive-condition (blank, nonsense) 2×2 evaluation matrix.

	Follow ($k_{\text{wrong}} \in \mathcal{A}(S_T)$)	Ignore ($k^* \in \mathcal{A}(S_T)$)	Fail (Neither Goal)
I_{orig} Succeeded	<i>Instruction-following.</i> Switched to newly instructed goal when asked.	<i>Scene-prior override.</i> Defaulted to original task despite new prompt.	<i>Confused.</i> Misleading prompt broke policy without redirection.
I_{orig} Failed	<i>Instruction-following.</i> Completed new goal despite failing baseline prompt.	<i>Noisy execution.</i> Achieved original goal unreliably.	<i>Uninformative.</i> Failed task under both conditions.

Table 3: The misleading-condition (wrong_object, wrong_action, wrong_task) 3×2 evaluation matrix.

LIBERO-Goal in MuJoCo. All tasks share an identical initial scene layout and object inventory, ensuring that the instruction is the sole distinguishing input.

5.2 Scene-Fixed Protocol

For each (task, episode) pair, we hold the initial simulator state S_0 strictly constant across all instruction conditions: $S_0(I_{\text{orig}}) \equiv S_0(I_{\text{pert}})$. We audit this invariant by hashing simulator state arrays (`qpos`, `qvel`) immediately after environment reset. As detailed in Section B, un-audited evaluations frequently violate this invariant due to simulator generator drift across process resets.

5.3 Data Collection and Implementation

Each rollout logs a structured record containing model checkpoint hash, task ID, condition, seed, initial-state hash, achieved task set $\mathcal{A}(S_T)$, and end-effector kinematics. Rollouts run for $T = 300$ steps under seeds 7 and 42. Simulator evaluations execute sequentially with one EGL rendering context per process to prevent rendering context crashes.

6 Results and Analysis

6.1 RQ1: Language is Causally Load-Bearing and Scale-Invariant

Table 4 presents Task Success Rate (TSR), Causal Sensitivity Score (CSS), and Original-Action Rate (OAR) under structural and semantic perturbations.

Destructive conditions (blank, nonsense) collapse task success across all architectures. SmolVLA-450M drops to 0.0% success (CSS = 1.00). OpenVLA-7B drops to 0.0% under blank

(CSS = 1.00). OpenVLA-OFT retains a minor residue (12.5% under blank, CSS = 0.87), which we attribute to visual shortcuts enabled by its wrist camera and proprioception inputs. Crucially, causal sensitivity is flat across model scale: CSS remains near 1.00 from 0.45B to 7.5B parameters.

6.2 RQ2: Information Structure Confounds Word-Class Grounding

Prior studies observe that object swaps collapse performance while verb swaps are tolerated, concluding that VLAs bind nouns strictly but ignore verbs. We find that this conclusion is confounded by the benchmark’s information structure.

We compute the conditional task entropy $H(\text{task} \mid W)$ over LIBERO-Goal’s ten task instructions under a uniform task prior ($H(\text{task}) = \log_2 10 = 3.32$ bits):

$$\begin{aligned} H(\text{task} \mid \text{noun set}) &= 0.00 \text{ bits}, \\ H(\text{task} \mid \text{verb set}) &= 1.55 \text{ bits}. \end{aligned} \quad (3)$$

All ten noun sets are unique across tasks, resolving all 3.32 bits of task identity. In contrast, the verb `put` appears in six of the ten tasks, leaving 1.55 bits unresolved (Figure 1). Conditional on knowing the nouns, the verb contributes *exactly zero bits* of task information.

This information redundancy has two major implications: 1. Ignoring the verb is informationally optimal on LIBERO-Goal. 2. Standard verb substitution probes are uninformative by construction: 7 of the 10 verb swaps replace `put` with `push`, leaving the noun set and task identity unchanged.

Model	Baseline	Blank	Nonsense	Wrong Action	Wrong Object	Wrong Task	Repeated
SmolVLA (0.45B)	78.5%±2.5	0.0% (1.00)	0.0% (1.00)	64.5%±0.5	9.3%±2.1	0.0%	10.0%±2.0
OpenVLA (7B)	71.0%±1.0	0.0% (1.00)	0.0% (1.00)	59.0%	10.7%±0.7	0.0%	58.5%±0.5
OpenVLA-OFT (7.5B)	97.5%±1.5	12.5%±2.5 (0.87)	14.0%±1.0 (0.86)	94.5%±0.5	11.4%	0.0%	96.5%±1.5
episodes per cell	200	200	200	200	140	200	200

Table 4: TSR under structural and semantic perturbations on LIBERO-Goal, with CSS in parentheses for destructive conditions. Episode counts differ by condition and model, so there is no single N : original $n = 200$, blank $n = 200$, nonsense $n = 200$, wrong_action $n = 200$, wrong_object $n = 140$, wrong_task $n = 200$, repeated $n = 200$. Seeds 42, 7. All models pass the scene-fixed check on every comparable episode.




Figure 1: LIBERO-Goal information structure. (a) Nouns resolve all 3.32 bits of task identity; verbs leave 1.55 bits unresolved. (b) Verbs partition six tasks into a single uninformative put cluster.

6.3 The Word-Class Removal Test

To test this entropy account, we perform a removal test: rather than substituting words, we drop word classes outright. Deleting the verb (0 bits destroyed) should leave performance unchanged, whereas masking nouns (3.32 bits destroyed) should collapse performance.

The empirical results confirm our prediction (Table 5, Figure 2). Deleting the verb changes success by only 2.0 points on OpenVLA-OFT (McNemar $p = 0.34$, 10 changed outcomes out of 200 matched pairs, consistent with random chance). Masking nouns costs 73.5 points and flips 148 pairs to failure. Two edits of similar edit distance differ by two orders of magnitude in impact because of their information content.

6.4 Re-evaluating Word-Class Asymmetry: Antonym vs. Near-Synonym Split

We separate the ten verb substitutions into the 3 genuine antonym pairs (open↔close twice, turn on↔turn off once) and the 7 near-synonym control pairs (Table 6, Figure 3).

Splitting the probe reverses the conclusion for OpenVLA-7B. Evaluating the suite-wide

Condition	bits	TSR	Δ [95% CI]	p
<i>OpenVLA (7B)</i>				
original	0.00	71.0%	0.0	–
verb_dropped	0.00	75.0%	+4.0 [−3.5, +11.0]	0.36
nouns_masked	3.32	6.0%	−65.0 [−71.5, −58.0]	$< 10^{-10}$
blank	3.32	0.0%	−71.0 [−77.0, −64.5]	$< 10^{-10}$
<i>OpenVLA-OFT (7.5B)</i>				
original	0.00	97.5%	0.0	–
verb_dropped	0.00	95.5%	−2.0 [−5.0, +1.0]	0.34
nouns_masked	3.32	24.0%	−73.5 [−79.5, −67.5]	$< 10^{-10}$
blank	3.32	12.5%	−85.0 [−89.5, −80.0]	$< 10^{-10}$

Table 5: Word-class removal test ($n = 200$ per cell). Coverage: OpenVLA (7B) on 10 of 10 tasks; OpenVLA-OFT (7.5B) on 10 of 10 tasks. The ablation conditions are generated in the same process as original, so the Δ column is a within-scene contrast.




Figure 2: Removal test on OpenVLA-OFT. Deleting the verb costs 2.0 points ($p = 0.34$), while masking nouns costs 73.5 points.

wrong_action probe yields a mild 12.0-point drop. However, splitting the probe reveals that OpenVLA-7B drops 31.7 points on true antonyms ($p < 10^{-5}$), while dropping only 3.6 points on near-synonyms ($p = 0.52$). OpenVLA-7B actively reads and grounds action verbs when the prompt contains true task contrast; suite-wide probes mask this by averaging real effects with uninformative controls.

Model	antonym		near-synonym	
	Δ TSR	p	Δ TSR	p
SmolVLA (0.45B)	-21.7 pp	10^{-3}	-10.7 pp	0.02
OpenVLA (7B)	-31.7 pp	10^{-6}	-3.6 pp	0.52
OpenVLA-OFT (7.5B)	-5.0 pp	0.45	-2.1 pp	0.51

Table 6: TSR drop under verb substitution split by task identity change (p -value from McNemar test over matched pairs).




Figure 3: Verb substitution split into antonyms ($n = 60$) versus near-synonyms ($n = 140$). OpenVLA-7B drops 31.7 points on antonyms ($p < 10^{-5}$).

6.5 Paraphrase Robustness and Kinematics

Table 7 reports PRIDE paraphrase scores. Robustness scales with parameter count: SmolVLA scores PRIDE 2.7, OpenVLA-7B scores 33.3, and OpenVLA-OFT achieves 65.8. All models follow the same vulnerability ordering: object paraphrases degrade performance least, while compositional paraphrases degrade performance most.

Kinematic analysis (t_{div} , threshold 5cm) confirms these findings: under blank instructions, trajectories depart from canonical paths within 1–2 steps. Under `wrong_action`, trajectories follow canonical paths for 6–11 steps before diverging, matching the timing expected when nouns specify target approach location and verbs specify

Model	Paraphrase Axis	n	scenes	TSR	Δ TSR (pp)	PRIDE
<i>SmolVLA (0.45B)</i>						
SmolVLA (0.45B)	Object	259	20	29.73%	-48.77	28.7
SmolVLA (0.45B)	Action	870	20	7.24%	-71.26	4.8
SmolVLA (0.45B)	Compositional	2963	20	1.79%	-76.71	1.5
SmolVLA (0.45B)	All axes	4092	24	4.72%	-73.78	2.7
<i>OpenVLA (7B)</i>						
OpenVLA (7B)	Object	150	150	48.00%	-23.00	36.3
OpenVLA (7B)	Action	150	150	60.67%	-10.33	51.8
OpenVLA (7B)	Compositional	150	150	24.00%	-47.00	20.2
OpenVLA (7B)	All axes	450	150	44.22%	-26.78	33.3
<i>OpenVLA-OFT (7.5B)</i>						
OpenVLA-OFT (7.5B)	Object	150	150	78.00%	-19.50	72.5
OpenVLA-OFT (7.5B)	Action	150	150	90.67%	-6.83	86.3
OpenVLA-OFT (7.5B)	Compositional	150	150	54.00%	-43.50	50.3
OpenVLA-OFT (7.5B)	All axes	450	150	74.22%	-23.28	65.8

Table 7: Paraphrase robustness scored via PRIDE.

manipulation manner.

7 Discussion and Conclusion

We evaluated three VLA architectures under a scene-fixed causal protocol to determine the nature of language grounding failures. Our findings establish three key conclusions:

1. ****Language is causally necessary:**** Completely removing language collapses task success across all scales, disproving the notion that VLAs operate as purely visual policies on shared-scene benchmarks.
2. ****Benchmark information redundancy creates apparent verb blindness:**** In LIBERO-Goal, nouns uniquely determine task identity. Models ignore verbs because verbs are redundant, not because models are incapable of verb grounding.
3. ****VLAs ground verbs when prompts are contrastive:**** On true verb antonyms, OpenVLA-7B exhibits strong verb sensitivity (31.7-point drop, $p < 10^{-5}$).

We recommend that future VLA benchmarks measure and report instruction set information content before drawing word-class grounding conclusions.

Limitations

Our study has several scope bounds: (1) **Single benchmark suite:** Causal isolation relies on LIBERO-Goal’s shared initial scene. (2) **Antonym task coverage:** Only three tasks in LIBERO-Goal afford true verb antonyms ($n = 60$ episodes across 3 task templates). (3) **Simulation scope:** All experiments take place in MuJoCo simulation without real-robot deployment. (4) **Uniform entropy prior:** Information calculations assume a uniform task distribution, which may differ from model pre-training data distributions.

Ethics and Reproducibility

This work uses simulated environments and open-source model checkpoints. We release per-run manifests (seed, environment locks, state hashes) to ensure full computational reproducibility.

References

Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, Pete Florence, Chuyuan Fu, Montse Gonzalez Arenas, Keerthana Gopalakrishnan, Kehang Han,

489	Karol Hausman, Alex Herzog, Jasmine Hsu, Brian	Chanyoung Kim, Minwoo Kim, Minseok Kang, Hyun-	547
490	Ichter, and 35 others. 2023. Rt-2: Vision-language-	woo Kim, and Dahuin Jung. 2026. Libero-para: A	548
491	action models transfer web knowledge to robotic	diagnostic benchmark and metrics for paraphrase ro-	549
492	control. In <i>arXiv preprint arXiv:2307.15818</i> .	bustness in vla models . <i>Preprint</i> , arXiv:2603.28301.	550
493	Anthony Brohan, Noah Brown, Justice Carbajal,	Moo Jin Kim, Chelsea Finn, and Percy Liang.	551
494	Yevgen Chebotar, Joseph Dabis, Chelsea Finn,	2025a. Fine-tuning vision-language-action mod-	552
495	Keerthana Gopalakrishnan, Karol Hausman, Alex	els: Optimizing speed and success. <i>arXiv preprint</i>	553
496	Herzog, Jasmine Hsu, Julian Ibarz, Brian Ichter,	<i>arXiv:2502.19645</i> .	554
497	Alex Irpan, Tomas Jackson, Sally Jesmonth, Nikhil	Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti,	555
498	Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng	Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael	556
499	Kuang, and 32 others. 2022. Rt-1: Robotics trans-	Rafailov, Ethan P Foster, Pannag R Sanketi, Quan	557
500	former for real-world control at scale. In <i>arXiv</i>	Vuong, Thomas Kollar, Benjamin Burchfiel, Russ	558
501	<i>preprint arXiv:2212.06817</i> .	Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang,	559
502	Open X-Embodiment Collaboration, Abby O’Neill,	and Chelsea Finn. 2025b. Openvla: An open-source	560
503	Abdul Rehman, Abhinav Gupta, Abhiram Mad-	vision-language-action model. In <i>Proceedings of</i>	561
504	dukuri, Abhishek Gupta, Abhishek Padalkar, Abra-	<i>The 8th Conference on Robot Learning</i> .	562
505	ham Lee, Acorn Pooley, Agrim Gupta, Ajay Man-	Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang	563
506	dlekar, Ajinkya Jain, Albert Tung, Alex Bewley,	Liu, Yuke Zhu, and Peter Stone. 2023. Libero:	564
507	Alex Herzog, Alex Irpan, Alexander Khazatsky,	Benchmarking knowledge transfer for lifelong robot	565
508	Anant Rai, Anchit Gupta, and 275 others. 2023.	learning. <i>arXiv preprint arXiv:2306.03310</i> .	566
509	Open X-Embodiment: Robotic learning datasets	Tom McCoy, Ellie Pavlick, and Tal Linzen. 2019.	567
510	and RT-X models. https://arxiv.org/abs/	Right for the wrong reasons: Diagnosing syntactic	568
511	2310.08864 .	heuristics in natural language inference. In <i>Proceed-</i>	569
512	Danny Driess, Fei Xia, Mehdi S. M. Sajjadi, Corey	<i>ings of the 57th Annual Meeting of the Association</i>	570
513	Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan	<i>for Computational Linguistics</i> , pages 3428–3448.	571
514	Wahid, Jonathan Tompson, Quan Vuong, Tianhe	Yiran Pang, Yiheng Zhao, Zhuopu Zhou, Tingkai Hu,	572
515	Yu, Wenlong Huang, Yevgen Chebotar, Pierre Ser-	and Ranxin Hou. 2025. Is OpenVLA truly robust?	573
516	manet, Daniel Duckworth, Sergey Levine, Vincent	a systematic evaluation of positional robustness. In	574
517	Vanhoucke, Karol Hausman, Marc Toussaint, Klaus	<i>Proceedings of IJCNLP-AACL</i> .	575
518	Greff, and 3 others. 2023. Palm-e: An embod-	Mustafa Shukor, Dana Aubakirova, Francesco Ca-	576
519	ied multimodal language model. In <i>arXiv preprint</i>	puano, Pepijn Kooijmans, Steven Palma, Adil	577
520	<i>arXiv:2303.03378</i> .	Zouitine, Michel Aractingi, Caroline Pascal, Mar-	578
521	Senyu Fei, Siyin Wang, Junhao Shi, Zihao Dai, Jikun	tino Russi, Andres Marafioti, Simon Alibert,	579
522	Cai, Pengfang Qian, Li Ji, Xinzhe He, Shiduo Zhang,	Matthieu Cord, Thomas Wolf, and Remi Cadene.	580
523	Zhaoye Fei, Jinlan Fu, Jingjing Gong, and Xipeng	2025. Smolvla: A vision-language-action model	581
524	Qiu. 2025. Libero-plus: In-depth robustness anal-	for affordable and efficient robotics . <i>Preprint</i> ,	582
525	ysis of vision-language-action models . <i>Preprint</i> ,	arXiv:2506.01844.	583
526	arXiv:2510.13626.	Eric Zhan and Kevin Wang. 2026. Evaluating sta-	584
527	Yash Goyal Goyal, Tushar Khot, Aishwarya Agrawal,	ble task execution in robot learning . <i>Preprint</i> ,	585
528	Douglas Zhang, Devi Parikh, and Dhruv Batra. 2017.	arXiv:2601.00100.	586
529	Making the v in vqa matter: Elevating the role of im-	Kevin Zhang and Serena Li. 2026. Loss landscape flat-	587
530	age understanding in visual question answering. In	ness and policy robustness in vision-language-action	588
531	<i>Proceedings of the IEEE Conference on Computer</i>	models . <i>Preprint</i> , arXiv:2601.00200.	589
532	<i>Vision and Pattern Recognition</i> , pages 6904–6913.	Xueyang Zhou, Yangming Xu, Guiyao Tie, Yongchao	590
533	Bryce Grant, Xijia Zhao, and Peng Wang. 2026.	Chen, Guowen Zhang, Duanfeng Chu, Pan Zhou,	591
534	Not all features are created equal: A mechanistic	and Lichao Sun. 2026. Libero-pro: Towards robust	592
535	study of vision-language-action models . <i>Preprint</i> ,	and fair evaluation of vision-language-action models	593
536	arXiv:2603.19233.	beyond memorization . <i>Preprint</i> , arXiv:2510.03827.	594
537	Suchin Gururangan, Swabha Swayamdipta, Omer	A Full Per-Cell Results	595
538	Levy, Roy Roy, Samuel R Bowman, and Noah A	Table 8 provides per-cell episode counts and 95%	596
539	Smith. 2018. Annotation artifacts in natural lan-	Wilson confidence intervals for all RQ1 conditions.	597
540	guage inference data. In <i>Proceedings of the 2018</i>	Table 9 reports kinematic divergence metrics.	598
541	<i>Conference of the North American Chapter of the</i>		
542	<i>Association for Computational Linguistics: Human</i>		
543	<i>Language Technologies</i> , pages 107–112.		
544	Yuchen Hou and Lin Zhao. 2026. Langgap: Diagnos-		
545	ing and closing the language gap in vision-language-		
546	action models . <i>Preprint</i> , arXiv:2603.00592.		

Model	Condition	n	TSR (95% CI)	CSS
SmolVLA (0.45B)	original	200	78.5% [72.3, 83.6]	–
SmolVLA (0.45B)	blank	200	0.0% [0.0, 1.9]	1.00
SmolVLA (0.45B)	nonsense	200	0.0% [0.0, 1.9]	1.00
SmolVLA (0.45B)	wrong_action	200	64.5% [57.7, 70.8]	–
SmolVLA (0.45B)	wrong_object	140	9.3% [5.5, 15.2]	–
SmolVLA (0.45B)	wrong_task	200	0.0% [0.0, 1.9]	–
SmolVLA (0.45B)	repeated	200	10.0% [6.6, 14.9]	–
OpenVLA (7B)	original	200	71.0% [64.4, 76.8]	–
OpenVLA (7B)	blank	200	0.0% [0.0, 1.9]	1.00
OpenVLA (7B)	nonsense	200	0.0% [0.0, 1.9]	1.00
OpenVLA (7B)	wrong_action	200	59.0% [52.1, 65.6]	–
OpenVLA (7B)	wrong_object	140	10.7% [6.6, 16.9]	–
OpenVLA (7B)	wrong_task	200	0.0% [0.0, 1.9]	–
OpenVLA (7B)	repeated	200	58.5% [51.6, 65.1]	–
OpenVLA-OFT (7.5B)	original	200	97.5% [94.3, 98.9]	–
OpenVLA-OFT (7.5B)	blank	200	12.5% [8.6, 17.8]	0.87
OpenVLA-OFT (7.5B)	nonsense	200	14.0% [9.9, 19.5]	0.86
OpenVLA-OFT (7.5B)	wrong_action	200	94.5% [90.4, 96.9]	–
OpenVLA-OFT (7.5B)	wrong_object	140	11.4% [7.2, 17.8]	–
OpenVLA-OFT (7.5B)	wrong_task	200	0.0% [0.0, 1.9]	–
OpenVLA-OFT (7.5B)	repeated	200	96.5% [93.0, 98.3]	–

Table 8: Full RQ1 per-cell success rates with 95% Wilson confidence intervals.

B Scene-Fixed Invariant Audit

Standard LIBERO reset logic advances an internal generator across calls. As a result, sequential conditions run in separate processes silently receive different initial object poses. We enforce per-task process execution and validate state hashes (qp_{os} , qv_{el}) to guarantee strict scene matching across conditions.

C PRIDE Metric Formulation

PRIDE (Kim et al., 2026) weights paraphrase success by distance from canonical instructions: $PD_i = 1 - 0.5(SK_i + ST_i)$ and $PRIDE = 100 \sum_i s_i PD_i / \sum_i PD_i$.

D Kinematic Divergence Analysis

Figure 4 illustrates end-effector divergence $e(t)$ over time step t .

E Initial Pilot Experiments on SmolVLA

In preliminary experiments on SmolVLA-450M using a 4-task subset of LIBERO-Object, blank instructions dropped success from 100% (8/8) to 0% (0/8), while object swaps dropped success to 0% (0/8), confirming early directional sensitivity before expanding to the full 10-task LIBERO-Goal suite.




Figure 4: End-effector displacement over rollout timesteps.

Model	Condition	pairs	t_{div} (samples)	e_{10} [m]
SmolVLA (0.45B)	blank	84	1.00	0.3548
SmolVLA (0.45B)	nonsense	84	1.87	0.2197
SmolVLA (0.45B)	para_action	10	2.00	0.1972
SmolVLA (0.45B)	para_compositional	13	1.85	0.2086
SmolVLA (0.45B)	para_object	13	4.69	0.1383
SmolVLA (0.45B)	repeated	84	2.40	0.1737
SmolVLA (0.45B)	wrong_action	77	9.31	0.0838
SmolVLA (0.45B)	wrong_task	84	1.64	0.3278
SmolVLA (0.45B)	wrong_object	56	1.96	0.1785
OpenVLA (7B)	blank	73	2.16	0.2315
OpenVLA (7B)	nonsense	73	1.92	0.2087
OpenVLA (7B)	nouns_masked	73	3.11	0.2015
OpenVLA (7B)	repeated	65	6.35	0.1082
OpenVLA (7B)	verb_dropped	68	7.87	0.0764
OpenVLA (7B)	wrong_action	68	6.18	0.1355
OpenVLA (7B)	wrong_task	73	2.25	0.2992
OpenVLA (7B)	wrong_object	48	2.75	0.1920
OpenVLA-OFT (7.5B)	blank	55	1.58	0.3236
OpenVLA-OFT (7.5B)	nonsense	54	1.65	0.3043
OpenVLA-OFT (7.5B)	nouns_masked	55	2.67	0.2751
OpenVLA-OFT (7.5B)	para_action	50	11.70	0.0339
OpenVLA-OFT (7.5B)	para_compositional	55	6.58	0.1474
OpenVLA-OFT (7.5B)	para_object	51	10.90	0.0582
OpenVLA-OFT (7.5B)	repeated	43	12.02	0.0234
OpenVLA-OFT (7.5B)	verb_dropped	53	12.13	0.0206
OpenVLA-OFT (7.5B)	wrong_action	51	11.37	0.0289
OpenVLA-OFT (7.5B)	wrong_task	55	1.64	0.3127
OpenVLA-OFT (7.5B)	wrong_object	35	1.89	0.3173

Table 9: Kinematic divergence step t_{div} and end-effector displacement e_{10} .